

Age, Period, and Cohort Effects in Philosophy Journal Citations

Anon

2025-02-12

There are extremely strong age and period effects in citations in philosophy journals. The age effect is that citations are concentrated on articles published two to five years prior. The period effect is that recent years have seen an explosion in the number of articles published, and the number of citations per articles, so many articles are getting more citations per year than they ever had previously. But cohort effects are trickier to detect. In this note I argue that they exist. There are more citations to articles from eras of more dramatic change in philosophy, such as around 1970 and around 2010. And there are fewer citations to articles from periods of consolidation, especially in the late 1970s through the 1980s.

1 Introduction

Before looking at the data, here are two things I believed about philosophy citations. First, philosophers tend to cite very old papers. We still regularly teach a number of papers over half a century old in introductory classes; e.g., (?), Thomson (1971), Singer (1972), Lewis (1973). These aren't taught as history papers, but as early entries into the contemporary philosophical debate. And, I thought, that's how we cite. Second, the technological changes of the last quarter century meant that this practice was being slowly reversed. The spread of electronic communication in the late 20th century, and then the rise of archives (e.g., Arxiv, SSRN, PhilPapers) and eventually journals publishing in EarlyView, meant that papers could now be cited even before they were published, and certainly without the delays involved in printing and posting journals around the world.

Both of these thoughts were wrong. Historically, philosophy papers have tended, when they are citing other philosophy papers, to cite very recent ones. But this tendency is diminishing, not increasing, over time. I'll offer much more evidence for these claims as we go along, but to make them plausible, I'll start with two simple graphs.

The data for the graphs come from citation data I downloaded concerning XXX papers published from 1955-2021,¹ in one hundred leading philosophy journals. I focussed on the citations to and from journals in this dataset. So every citation is from one of these 100 journals between 1955 and 2021, and to one of these 100 journals between 1955 and 2021. (The details of the journals, including when they start getting indexed for this dataset, are in [?@sec-methodology](#).) In total, that gives us YYY citations.

Say the *age* of a citation is the difference between the publication year of the citing article and the cited article. So if an article published in 1998 cites an article published in 1985, that's a 13 year old citation.

In Figure 1 I've plotted the number of citations in the dataset with each possible age. As you can see, it's very heavily tilted towards the left-hand edge. It is true that people still cite (?). Indeed, it's one of the most cited papers in the last ten years. But it's just one paper; the bulk of citations are to recently published papers which, if history is any guide, will soon stop collecting citations.

In Figure 2 I've plotted the median and mode age of citations in each year from 1980 onwards. Before that the numbers are even lower, but since I'm only looking at citations to articles published after 1955 (or later if Web of Science started indexing the journal later than that), this is arguably an artifact of how I'm collecting the data. From 1980 onwards, however, there are many older articles that could be, but are not, getting cited. The upwards trends in both graphs look like a real change in citation practices, and not in the direction I antecedently expected.

There is a third surprise in the data, but it's a little more equivocal, and I'm not sure what to make of

¹I would like to have more recent data, but this is the latest full year of data available through my university's contract with Web of Science. I do have substantial partial data for 2022, and it mostly confirms the trends shown here. But in this case I think it's better to leave off partial data than to try to correct for its incompleteness.

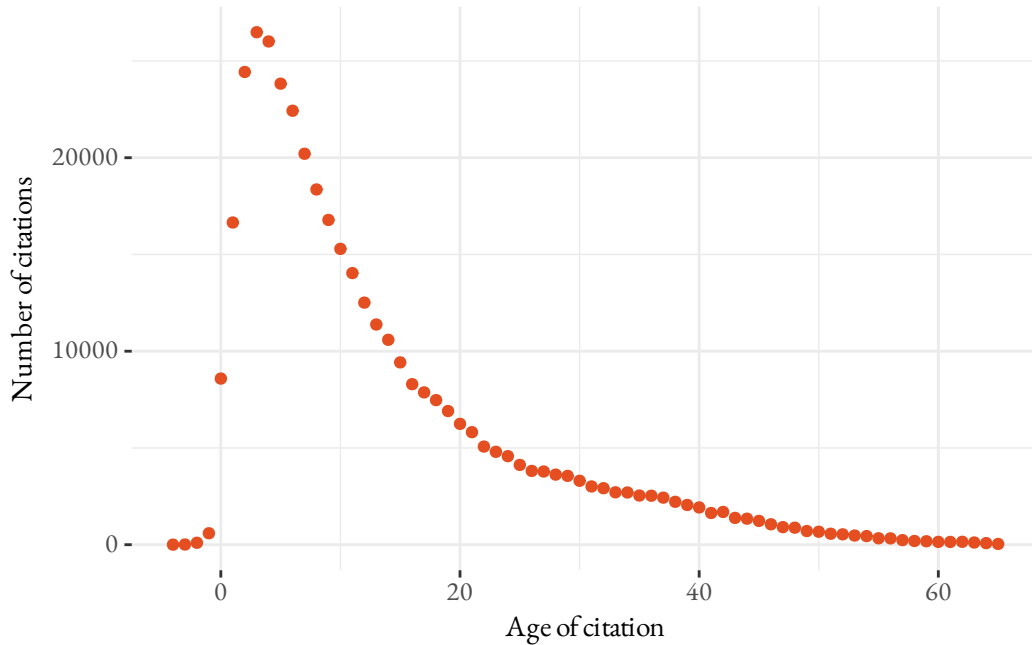




Figure 1: Number of citations with each age.

it. The 2010s seemed like, and to be honest still seem like, something of a golden age for philosophy. In metaphysics we saw the biggest paradigm shift in many years, away from modality and towards grounding. We saw the growth of important fields of social philosophy, including social epistemology, social metaphysics, and social philosophy of language. Though there were some earlier papers that have become important to the latter two fields (e.g., Haslanger (2000), and Langton (1993)), it would have been a stretch to even call them ‘fields’ before 2010. Social epistemology was always a bit bigger, and you could point to earlier field defining work by, e.g., Jennifer Lackey (2008) and Adam Elga (2007). But it grew phenomenally in the 2010s. I’d predicted that would show up in higher citations to work in the 2010s, as these changes were consolidated. The data are a bit messy, and it would be good to have much more data, but this does not look to have happened. There isn’t as neat a graph for this, however, and we’ll return to this point at the end.



(a) Median



(b) Mode

Figure 2: Summary statistics for outbound citations each year 1970-2021.

2 Age, Period, and Cohort

To help understand the citation patterns, I'll borrow some terminology that's common in both sociology and medicine. Imagine that we see, in the historical record, some interesting patterns among teenagers in the late 1960s, and we're wondering what could explain the pattern. Two types of pattern spring immediately to mind, along with ways to test them.

First, the behaviour could be explained by the fact the people involved are teenagers. If so, it is an **age effect**. The natural way to test this is to see if similar patterns show up with teenagers at different times.

Second, the behaviour could be explained by the fact that it was the 1960s, and lots of striking things happened in the 1960s. If so, it is a **period effect**. The natural way to test this is to see if the same pattern shows up with non-teenagers in the 1960s.

There is an important third kind of explanation. The people involved are born in the early 1950s, so they are part of the post-war baby boom. Colloquially, they are boomers. Maybe that could explain the pattern we see. If so, it is a **cohort effect**. The natural way to test this is to see if the same pattern shows up if we look at the same people in other stages of their life.

It's easy to overlook the importance of cohort effects. Sometimes they simply look like age effects. Ghitza, Gelman, and Auerbach (2023) argue that many hypotheses about age effects on voting, e.g., that older people are more naturally conservative, are really just cohort effects. Bump (2023) argues that understanding the distinctive role the boomers in particular play is crucial for understanding many aspects of modern American life.

There are mathematical reasons that it is hard to tease these effects apart too. Many statistical techniques for separating out influences start to fall apart when there are linear correlations between combinations of variables. In this case there is as tight a correlation as is possible. By definition, cohort plus age equals period. There are some things you can do to get around this problem - see Keyes et al. (2010) for a useful survey of some of the options - but it remains a challenge.

Even conceptually, it is hard to separate out these three effects in cases where there is evidence that the strength of the effects changes over time. As I noted at the start, the natural way to test hypotheses about which effect is strongest involve looking at other times. That works well when the age effects are constant. When they are not (and they might not be here), it is harder.

For most of our story, however, it helps just to have these three effects in mind. Using them, we can summarise the data reasonably quickly.

- The age effect is that articles get cited most when they are two to five years old, as shown in Figure 1.
- The period effect is that there are many more citations in recent years than in earlier years. This is in part because the number of articles published in these journals has been growing, and in part because the number of citations per article grew substantially over the 2000s and 2010s, and exploded in the 2020s.
- The cohort effect is that articles from the 1970s and 2000s get cited more than you'd expect given these age and period effects, articles from before the late-1960s get barely cited at all, and articles from 1980 through the mid-1990s also get cited considerably less than articles either side of that period. I'll offer some speculations at the end of the paper about the philosophical causes of, and consequences of, these cohort effects.

As I mentioned above, I'll go over the methodology in detail in [?@sec-methodology](#). But there is one point that is important to note before we start. I'm using data from Web of Science, and they typically don't start indexing journals until well after the journal is established. So the first year of citation data I have for *Analysis* is 1975. Crucially, that means that "Is Knowledge Justified True Belief?" (Gettier 1963) is not included in this study. If it were, and in general if I had the data from *Analysis* to work from, some of the results about the early 1960s would look less dramatic, though as far as I can tell, the direction of the results wouldn't change.

3 Period Effects

Those 397459 citations are not distributed evenly over time. Instead, they grow rapidly. At the start, in 1956, there are only 5 citations. That's not too surprising; without the ability to cite preprints, there aren't going to be many citations of articles that have come out that year. By 2021, there are 56300. In Figure 3, I show how these grew; the striking thing to me is the big jump between 2020 and 2021.




Figure 3: The number of citations in the dataset made each year.

What explains this dramatic growth? Part of the explanation is that more articles are being published, and more articles are being indexed. Figure 4 shows how many articles are in the dataset each year.

That explains some of the growth, but not all of it. The curve in Figure 4 is not nearly as steep as the curve in Figure 3. The number of (indexed) citations per article is also rising. In Figure 5 I've plotted the average number of citations to other articles in the dataset each year.

There are a few possible explanations for the shape of this graph.

At the left-hand edge, there are obvious boundary effects. Since we're only counting citations to articles




Figure 4: The number of articles in the dataset published each year.




Figure 5: The average number of citations to indexed articles each year.

published since 1956, it isn't surprising that there aren't very many of them per article in the 1950s. Since articles rarely get unpublished, there are more articles available to cite every year.

That can't explain the massive jumps we see at the right hand edge of Figure 5. The jump there looks like the convergence of two cultural trends. One is a trend simply to greater numbers of citations. The most casual perusal of journals will confirm that trend. The other is a trend to greater citations of journals themselves, as opposed to books or edited volumes.

A sharp jump like this is a warning sign that there is something wrong with the data. It's impractical to cross-check every entry, but those I have checked look correct. The change seems led by the most prestigious journals. For each journal I calculated the average number of outbound citations (to these hundred journal) for both the 2010s, and the first two years of the 2020s. The ten journals with the largest increase between the decades are shown in Table 1.

Table 1: Mean outbound citations for some journals over the last two decades.

Journal	2010-2019	2020-2021	Difference
Philosophical Review	14.8	26.4	11.6
Philosophical Perspectives	11.3	19.6	8.3
Philosophy and Phenomenological Research	9.6	15.2	5.6
Journal of Philosophy	9.0	13.7	4.7
Philosophical Studies	9.0	13.6	4.6
Noûs	11.5	16.0	4.5
Philosophical Quarterly	8.8	13.3	4.5
Philosophy	4.0	8.3	4.3
Philosophy Compass	11.2	15.4	4.2
Ethics	8.4	12.3	3.8


Since *Philosophical Review* only publishes 10 to 12 articles per year, it is not surprising that it shows the most variation on this list. Still, the change in the 2010s isn't only small sample size variation. Of the 22 articles it published in 2020 and 2021, only one of them (?) had fewer than 14.8 outbound citations. With a sample of just 22 anything could happen, but it would be surprising to have all but one end up on the same side of the historical average by chance.

Although the number of citations is going up, the number of articles available to be cited is also going up. Say an article is *available* if it is published in a year iff it is published in or before that year. That's not quite right in either direction; some articles are cited before publication, some articles that come out in December aren't in any real sense available to be cited in January. But it's close enough. Say an article is from a year that is *typically* cited iff it is between 3 and 10 years before the citing year. This notion will play a big role in ?@sec-age; I'm going to use these as a way of getting something like a base rate for citations in a given year. Using these definitions, Figure 6 shows how many articles are available to be cited each year, and are from years that are typically cited.


In Figure 7, I've shown how often, in each year, the available articles, and the 'typical' articles are cited. The 'available' graph is obviously similar to Figure 3; under 1% of citations are to articles published in future years. One thing that will be useful in ?@sec-age is that the graphs in Figure 7 have a similar shape.

Putting all these together we can work out how often, on average, available articles, and typical articles, are cited in each year. The results are in Figure 8.

Three things stand out about Figure 8. One is that the two graphs have pretty similar shapes. Using citations from 3 to 10 years prior to the citing year is a pretty good proxy for all citations, and it turns out to be stable in other ways. A second is that both graphs are fairly flat for a long time. Between the mid 1970s and early 2000s they bounce around without moving much. Then they take off, and go through the roof in 2021. The other thing is that these are low numbers. For most of this study, an arbitrary article in one of these hundred journals was cited in one of those journals once a *decade*. Actually, since




(a) Available articles




(b) Typically cited articles

Figure 6: Article counts.




(a) Citations to available articles




(b) Citations to typical articles

Figure 7: Citation counts.



(a) Available articles



(b) Typical articles

Figure 8: Mean annual citations to different article kinds.

citation rates are extremely long-tailed, and mean rates are well above medians, that somewhat overstates how often the ‘average article’ was being cited. Frequent citation is very much not the norm.²

The various period effects are substantial; to get an reliable picture of the trends in citation patterns, we’re going to have to allow for them.

Bump, Philip. 2023. *The Aftermath: The Last Days of the Baby Boom and the Future of Power in America*. New York: Penguin Random House.

Elga, Adam. 2007. “Reflection and Disagreement.” *Noûs* 41 (3): 478–502. <https://doi.org/10.1111/j.1468-0068.2007.00656.x>.

Gettier, Edmund L. 1963. “Is Justified True Belief Knowledge?” *Analysis* 23 (6): 121–23. <https://doi.org/10.2307/3326922>.

Ghitza, Yair, Andrew Gelman, and Jonathan Auerbach. 2023. “The Great Society, Reagan’s Revolution, and Generations of Presidential Voting.” *American Journal of Political Science* 67 (3): 520–37. <https://doi.org/https://doi.org/10.1111/ajps.12713>.

Haslanger, Sally. 2000. “Gender and Race: (What) Are They? (What) Do We Want Them to Be?” *Noûs* 34 (1): 31–55. <https://doi.org/10.1111/0029-4624.00201>.

Keyes, Katherine M., Rebecca L. Utz, Whitney Robinson, and Guohua Li. 2010. “What Is a Cohort Effect? Comparison of Three Statistical Methods for Modeling Cohort Effects in Obesity Prevalence in the United States, 1971–2006.” *Social Science & Medicine* 70 (7): 1100–1108. <https://doi.org/10.1016/j.socscimed.2009.12.018>.

Lackey, Jennifer. 2008. *Learning from Words: Testimony as a Source of Knowledge*. Oxford: Oxford University Press.

Langton, Rae. 1993. “Speech Acts and Unspeakable Acts.” *Philosophy & Public Affairs* 22 (4): 293–330.

Lewis, David. 1973. “Causation.” *Journal of Philosophy* 70 (17): 556–67. <https://doi.org/10.2307/2025310>.

²In the long run the average number of times an article is cited equals the average number of citations per article. So it shouldn’t be too surprising that most article have just a handful of citations in philosophy journals.

Singer, Peter. 1972. "Famine, Affluence, and Morality." *Philosophy & Public Affairs* 1 (3): 229–43.

Thomson, Judith Jarvis. 1971. "A Defense of Abortion." *Philosophy and Public Affairs* 1 (1): 47–66.